

USER GUIDE

The User Guide covers all of pandas by topic area. Each of the subsections introduces a topic (such as “working with missing data”), and discusses how pandas approaches the problem, with many examples throughout.

Users brand-new to pandas should start with 10min.

For a high level summary of the pandas fundamentals, see *Intro to data structures* and *Essential basic functionality*.

Further information on any specific method can be obtained in the *API reference*. {{ header }}

2.1 10 minutes to pandas

This is a short introduction to pandas, geared mainly for new users. You can see more complex recipes in the *Cookbook*.

Customarily, we import as follows:

```
In [1]: import numpy as np
In [2]: import pandas as pd
```

2.1.1 Object creation

See the *Data Structure Intro section*.

Creating a Series by passing a list of values, letting pandas create a default integer index:

```
In [3]: s = pd.Series([1, 3, 5, np.nan, 6, 8])
In [4]: s
Out[4]:
0    1.0
1    3.0
2    5.0
3    NaN
4    6.0
5    8.0
dtype: float64
```

Creating a DataFrame by passing a NumPy array, with a datetime index and labeled columns:

```
In [5]: dates = pd.date_range("20130101", periods=6)
In [6]: dates
```

(continues on next page)

(continued from previous page)

```

Out [6]:
DatetimeIndex(['2013-01-01', '2013-01-02', '2013-01-03', '2013-01-04',
              '2013-01-05', '2013-01-06'],
              dtype='datetime64[ns]', freq='D')

In [7]: df = pd.DataFrame(np.random.randn(6, 4), index=dates, columns=list("ABCD"))

In [8]: df
Out [8]:
           A          B          C          D
2013-01-01  1.712776  0.336509 -0.945427  0.449729
2013-01-02 -0.177498  1.487290 -0.458333 -0.721907
2013-01-03 -0.139201 -2.090971  0.975329 -0.565294
2013-01-04  0.248276 -1.001166  0.126691 -0.381556
2013-01-05  0.279357 -0.002987 -0.566766  1.929640
2013-01-06 -1.161636 -0.594994  0.722651  0.895291

```

Creating a DataFrame by passing a dict of objects that can be converted to series-like.

```

In [9]: df2 = pd.DataFrame(
...:     {
...:         "A": 1.0,
...:         "B": pd.Timestamp("20130102"),
...:         "C": pd.Series(1, index=list(range(4)), dtype="float32"),
...:         "D": np.array([3] * 4, dtype="int32"),
...:         "E": pd.Categorical(["test", "train", "test", "train"]),
...:         "F": "foo",
...:     }
...: )
...:

In [10]: df2
Out [10]:
           A          B          C          D          E          F
0  1.0 2013-01-02  1.0  3  test  foo
1  1.0 2013-01-02  1.0  3  train foo
2  1.0 2013-01-02  1.0  3  test  foo
3  1.0 2013-01-02  1.0  3  train foo

```

The columns of the resulting DataFrame have different *dtypes*.

```

In [11]: df2.dtypes
Out [11]:
A          float64
B  datetime64[ns]
C          float32
D           int32
E          category
F           object
dtype: object

```

If you're using IPython, tab completion for column names (as well as public attributes) is automatically enabled. Here's a subset of the attributes that will be completed:

```

In [12]: df2.<TAB> # noqa: E225, E999
df2.A          df2.bool

```

(continues on next page)

(continued from previous page)

df2.abs	df2.boxplot
df2.add	df2.C
df2.add_prefix	df2.clip
df2.add_suffix	df2.columns
df2.align	df2.copy
df2.all	df2.count
df2.any	df2.combine
df2.append	df2.D
df2.apply	df2.describe
df2.applymap	df2.diff
df2.B	df2.duplicated

As you can see, the columns A, B, C, and D are automatically tab completed. E and F are there as well; the rest of the attributes have been truncated for brevity.

2.1.2 Viewing data

See the *Basics section*.

Here is how to view the top and bottom rows of the frame:

```
In [13]: df.head()
Out[13]:
```

	A	B	C	D
2013-01-01	1.712776	0.336509	-0.945427	0.449729
2013-01-02	-0.177498	1.487290	-0.458333	-0.721907
2013-01-03	-0.139201	-2.090971	0.975329	-0.565294
2013-01-04	0.248276	-1.001166	0.126691	-0.381556
2013-01-05	0.279357	-0.002987	-0.566766	1.929640

```
In [14]: df.tail(3)
Out[14]:
```

	A	B	C	D
2013-01-04	0.248276	-1.001166	0.126691	-0.381556
2013-01-05	0.279357	-0.002987	-0.566766	1.929640
2013-01-06	-1.161636	-0.594994	0.722651	0.895291

Display the index, columns:

```
In [15]: df.index
Out[15]:
DatetimeIndex(['2013-01-01', '2013-01-02', '2013-01-03', '2013-01-04',
               '2013-01-05', '2013-01-06'],
              dtype='datetime64[ns]', freq='D')

In [16]: df.columns
Out[16]: Index(['A', 'B', 'C', 'D'], dtype='object')
```

`DataFrame.to_numpy()` gives a NumPy representation of the underlying data. Note that this can be an expensive operation when your DataFrame has columns with different data types, which comes down to a fundamental difference between pandas and NumPy: **NumPy arrays have one dtype for the entire array, while pandas DataFrames have one dtype per column.** When you call `DataFrame.to_numpy()`, pandas will find the NumPy dtype that can hold *all* of the dtypes in the DataFrame. This may end up being `object`, which requires casting every value to a Python object.

For `df`, our `DataFrame` of all floating-point values, `DataFrame.to_numpy()` is fast and doesn't require copying data.

```
In [17]: df.to_numpy()
Out[17]:
array([[ 1.71277599,  0.33650864, -0.94542653,  0.44972919],
       [-0.17749784,  1.48729025, -0.45833313, -0.72190726],
       [-0.1392012 , -2.09097127,  0.97532913, -0.5652943 ],
       [ 0.248276  , -1.00116557,  0.12669091, -0.38155642],
       [ 0.27935679, -0.00298719, -0.56676602,  1.92963962],
       [-1.16163559, -0.59499379,  0.72265057,  0.89529069]])
```

For `df2`, the `DataFrame` with multiple dtypes, `DataFrame.to_numpy()` is relatively expensive.

```
In [18]: df2.to_numpy()
Out[18]:
array([[1.0, Timestamp('2013-01-02 00:00:00'), 1.0, 3, 'test', 'foo'],
       [1.0, Timestamp('2013-01-02 00:00:00'), 1.0, 3, 'train', 'foo'],
       [1.0, Timestamp('2013-01-02 00:00:00'), 1.0, 3, 'test', 'foo'],
       [1.0, Timestamp('2013-01-02 00:00:00'), 1.0, 3, 'train', 'foo']],
      dtype=object)
```

Note: `DataFrame.to_numpy()` does *not* include the index or column labels in the output.

`describe()` shows a quick statistic summary of your data:

```
In [19]: df.describe()
Out[19]:
```

	A	B	C	D
count	6.000000	6.000000	6.000000	6.000000
mean	0.127012	-0.311053	-0.024309	0.267650
std	0.935604	1.222560	0.763036	1.027987
min	-1.161636	-2.090971	-0.945427	-0.721907
25%	-0.167924	-0.899623	-0.539658	-0.519360
50%	0.054537	-0.298990	-0.165821	0.034086
75%	0.271587	0.251635	0.573661	0.783900
max	1.712776	1.487290	0.975329	1.929640

Transposing your data:

```
In [20]: df.T
Out[20]:
```

	2013-01-01	2013-01-02	2013-01-03	2013-01-04	2013-01-05	2013-01-06
A	1.712776	-0.177498	-0.139201	0.248276	0.279357	-1.161636
B	0.336509	1.487290	-2.090971	-1.001166	-0.002987	-0.594994
C	-0.945427	-0.458333	0.975329	0.126691	-0.566766	0.722651
D	0.449729	-0.721907	-0.565294	-0.381556	1.929640	0.895291

Sorting by an axis:

```
In [21]: df.sort_index(axis=1, ascending=False)
Out[21]:
```

	D	C	B	A
2013-01-01	0.449729	-0.945427	0.336509	1.712776
2013-01-02	-0.721907	-0.458333	1.487290	-0.177498
2013-01-03	-0.565294	0.975329	-2.090971	-0.139201

(continues on next page)

(continued from previous page)

```

2013-01-04 -0.381556  0.126691 -1.001166  0.248276
2013-01-05  1.929640 -0.566766 -0.002987  0.279357
2013-01-06  0.895291  0.722651 -0.594994 -1.161636

```

Sorting by values:

```

In [22]: df.sort_values(by="B")
Out[22]:
           A          B          C          D
2013-01-03 -0.139201 -2.090971  0.975329 -0.565294
2013-01-04  0.248276 -1.001166  0.126691 -0.381556
2013-01-06 -1.161636 -0.594994  0.722651  0.895291
2013-01-05  0.279357 -0.002987 -0.566766  1.929640
2013-01-01  1.712776  0.336509 -0.945427  0.449729
2013-01-02 -0.177498  1.487290 -0.458333 -0.721907

```

2.1.3 Selection

Note: While standard Python / NumPy expressions for selecting and setting are intuitive and come in handy for interactive work, for production code, we recommend the optimized pandas data access methods, `.at`, `.iat`, `.loc` and `.iloc`.

See the indexing documentation [Indexing and Selecting Data](#) and [MultiIndex / Advanced Indexing](#).

Getting

Selecting a single column, which yields a `Series`, equivalent to `df.A`:

```

In [23]: df["A"]
Out[23]:
2013-01-01    1.712776
2013-01-02   -0.177498
2013-01-03   -0.139201
2013-01-04    0.248276
2013-01-05    0.279357
2013-01-06   -1.161636
Freq: D, Name: A, dtype: float64

```

Selecting via `[]`, which slices the rows.

```

In [24]: df[0:3]
Out[24]:
           A          B          C          D
2013-01-01  1.712776  0.336509 -0.945427  0.449729
2013-01-02 -0.177498  1.487290 -0.458333 -0.721907
2013-01-03 -0.139201 -2.090971  0.975329 -0.565294

In [25]: df["20130102":"20130104"]
Out[25]:
           A          B          C          D
2013-01-02 -0.177498  1.487290 -0.458333 -0.721907

```

(continues on next page)

(continued from previous page)

```
2013-01-03 -0.139201 -2.090971  0.975329 -0.565294
2013-01-04  0.248276 -1.001166  0.126691 -0.381556
```

Selection by label

See more in [Selection by Label](#).

For getting a cross section using a label:

```
In [26]: df.loc[dates[0]]
Out[26]:
A    1.712776
B     0.336509
C   -0.945427
D     0.449729
Name: 2013-01-01 00:00:00, dtype: float64
```

Selecting on a multi-axis by label:

```
In [27]: df.loc[:, ["A", "B"]]
Out[27]:
           A         B
2013-01-01  1.712776  0.336509
2013-01-02 -0.177498  1.487290
2013-01-03 -0.139201 -2.090971
2013-01-04  0.248276 -1.001166
2013-01-05  0.279357 -0.002987
2013-01-06 -1.161636 -0.594994
```

Showing label slicing, both endpoints are *included*:

```
In [28]: df.loc["20130102":"20130104", ["A", "B"]]
Out[28]:
           A         B
2013-01-02 -0.177498  1.487290
2013-01-03 -0.139201 -2.090971
2013-01-04  0.248276 -1.001166
```

Reduction in the dimensions of the returned object:

```
In [29]: df.loc["20130102", ["A", "B"]]
Out[29]:
A    -0.177498
B     1.487290
Name: 2013-01-02 00:00:00, dtype: float64
```

For getting a scalar value:

```
In [30]: df.loc[dates[0], "A"]
Out[30]: 1.7127759912633918
```

For getting fast access to a scalar (equivalent to the prior method):

```
In [31]: df.at[dates[0], "A"]
Out[31]: 1.7127759912633918
```

Selection by position

See more in *Selection by Position*.

Select via the position of the passed integers:

```
In [32]: df.iloc[3]
Out[32]:
A    0.248276
B   -1.001166
C    0.126691
D   -0.381556
Name: 2013-01-04 00:00:00, dtype: float64
```

By integer slices, acting similar to NumPy/Python:

```
In [33]: df.iloc[3:5, 0:2]
Out[33]:
              A          B
2013-01-04  0.248276 -1.001166
2013-01-05  0.279357 -0.002987
```

By lists of integer position locations, similar to the NumPy/Python style:

```
In [34]: df.iloc[[1, 2, 4], [0, 2]]
Out[34]:
              A          C
2013-01-02 -0.177498 -0.458333
2013-01-03 -0.139201  0.975329
2013-01-05  0.279357 -0.566766
```

For slicing rows explicitly:

```
In [35]: df.iloc[1:3, :]
Out[35]:
              A          B          C          D
2013-01-02 -0.177498  1.487290 -0.458333 -0.721907
2013-01-03 -0.139201 -2.090971  0.975329 -0.565294
```

For slicing columns explicitly:

```
In [36]: df.iloc[:, 1:3]
Out[36]:
              B          C
2013-01-01  0.336509 -0.945427
2013-01-02  1.487290 -0.458333
2013-01-03 -2.090971  0.975329
2013-01-04 -1.001166  0.126691
2013-01-05 -0.002987 -0.566766
2013-01-06 -0.594994  0.722651
```

For getting a value explicitly:

```
In [37]: df.iloc[1, 1]
Out[37]: 1.48729024797565
```

For getting fast access to a scalar (equivalent to the prior method):

```
In [38]: df.iat[1, 1]
Out[38]: 1.48729024797565
```

Boolean indexing

Using a single column's values to select data.

```
In [39]: df[df["A"] > 0]
Out[39]:
```

	A	B	C	D
2013-01-01	1.712776	0.336509	-0.945427	0.449729
2013-01-04	0.248276	-1.001166	0.126691	-0.381556
2013-01-05	0.279357	-0.002987	-0.566766	1.929640

Selecting values from a DataFrame where a boolean condition is met.

```
In [40]: df[df > 0]
Out[40]:
```

	A	B	C	D
2013-01-01	1.712776	0.336509	NaN	0.449729
2013-01-02	NaN	1.487290	NaN	NaN
2013-01-03	NaN	NaN	0.975329	NaN
2013-01-04	0.248276	NaN	0.126691	NaN
2013-01-05	0.279357	NaN	NaN	1.929640
2013-01-06	NaN	NaN	0.722651	0.895291

Using the `isin()` method for filtering:

```
In [41]: df2 = df.copy()
In [42]: df2["E"] = ["one", "one", "two", "three", "four", "three"]
In [43]: df2
Out[43]:
```

	A	B	C	D	E
2013-01-01	1.712776	0.336509	-0.945427	0.449729	one
2013-01-02	-0.177498	1.487290	-0.458333	-0.721907	one
2013-01-03	-0.139201	-2.090971	0.975329	-0.565294	two
2013-01-04	0.248276	-1.001166	0.126691	-0.381556	three
2013-01-05	0.279357	-0.002987	-0.566766	1.929640	four
2013-01-06	-1.161636	-0.594994	0.722651	0.895291	three

```
In [44]: df2[df2["E"].isin(["two", "four"])]
Out[44]:
```

	A	B	C	D	E
2013-01-03	-0.139201	-2.090971	0.975329	-0.565294	two
2013-01-05	0.279357	-0.002987	-0.566766	1.929640	four

Setting

Setting a new column automatically aligns the data by the indexes.

```
In [45]: s1 = pd.Series([1, 2, 3, 4, 5, 6], index=pd.date_range("20130102",
↳ periods=6))

In [46]: s1
Out[46]:
2013-01-02    1
2013-01-03    2
2013-01-04    3
2013-01-05    4
2013-01-06    5
2013-01-07    6
Freq: D, dtype: int64

In [47]: df["F"] = s1
```

Setting values by label:

```
In [48]: df.at[dates[0], "A"] = 0
```

Setting values by position:

```
In [49]: df.iat[0, 1] = 0
```

Setting by assigning with a NumPy array:

```
In [50]: df.loc[:, "D"] = np.array([5] * len(df))
```

The result of the prior setting operations.

```
In [51]: df
Out[51]:
```

	A	B	C	D	F
2013-01-01	0.000000	0.000000	-0.945427	5	NaN
2013-01-02	-0.177498	1.487290	-0.458333	5	1.0
2013-01-03	-0.139201	-2.090971	0.975329	5	2.0
2013-01-04	0.248276	-1.001166	0.126691	5	3.0
2013-01-05	0.279357	-0.002987	-0.566766	5	4.0
2013-01-06	-1.161636	-0.594994	0.722651	5	5.0

A where operation with setting.

```
In [52]: df2 = df.copy()

In [53]: df2[df2 > 0] = -df2

In [54]: df2
Out[54]:
```

	A	B	C	D	F
2013-01-01	0.000000	0.000000	-0.945427	-5	NaN
2013-01-02	-0.177498	-1.487290	-0.458333	-5	-1.0
2013-01-03	-0.139201	-2.090971	-0.975329	-5	-2.0
2013-01-04	-0.248276	-1.001166	-0.126691	-5	-3.0
2013-01-05	-0.279357	-0.002987	-0.566766	-5	-4.0
2013-01-06	-1.161636	-0.594994	-0.722651	-5	-5.0

2.1.4 Missing data

pandas primarily uses the value `np.nan` to represent missing data. It is by default not included in computations. See the [Missing Data section](#).

Reindexing allows you to change/add/delete the index on a specified axis. This returns a copy of the data.

```
In [55]: df1 = df.reindex(index=dates[0:4], columns=list(df.columns) + ["E"])
```

```
In [56]: df1.loc[dates[0] : dates[1], "E"] = 1
```

```
In [57]: df1
```

```
Out [57]:
```

	A	B	C	D	F	E
2013-01-01	0.000000	0.000000	-0.945427	5	NaN	1.0
2013-01-02	-0.177498	1.487290	-0.458333	5	1.0	1.0
2013-01-03	-0.139201	-2.090971	0.975329	5	2.0	NaN
2013-01-04	0.248276	-1.001166	0.126691	5	3.0	NaN

To drop any rows that have missing data.

```
In [58]: df1.dropna(how="any")
```

```
Out [58]:
```

	A	B	C	D	F	E
2013-01-02	-0.177498	1.48729	-0.458333	5	1.0	1.0

Filling missing data.

```
In [59]: df1.fillna(value=5)
```

```
Out [59]:
```

	A	B	C	D	F	E
2013-01-01	0.000000	0.000000	-0.945427	5	5.0	1.0
2013-01-02	-0.177498	1.487290	-0.458333	5	1.0	1.0
2013-01-03	-0.139201	-2.090971	0.975329	5	2.0	5.0
2013-01-04	0.248276	-1.001166	0.126691	5	3.0	5.0

To get the boolean mask where values are nan.

```
In [60]: pd.isna(df1)
```

```
Out [60]:
```

	A	B	C	D	F	E
2013-01-01	False	False	False	False	True	False
2013-01-02	False	False	False	False	False	False
2013-01-03	False	False	False	False	False	True
2013-01-04	False	False	False	False	False	True

2.1.5 Operations

See the [Basic section on Binary Ops](#).

Stats

Operations in general *exclude* missing data.

Performing a descriptive statistic:

```
In [61]: df.mean()
Out[61]:
A    -0.158450
B    -0.367138
C    -0.024309
D     5.000000
F     3.000000
dtype: float64
```

Same operation on the other axis:

```
In [62]: df.mean(1)
Out[62]:
2013-01-01    1.013643
2013-01-02    1.370292
2013-01-03    1.149031
2013-01-04    1.474760
2013-01-05    1.741921
2013-01-06    1.793204
Freq: D, dtype: float64
```

Operating with objects that have different dimensionality and need alignment. In addition, pandas automatically broadcasts along the specified dimension.

```
In [63]: s = pd.Series([1, 3, 5, np.nan, 6, 8], index=dates).shift(2)
In [64]: s
Out[64]:
2013-01-01    NaN
2013-01-02    NaN
2013-01-03     1.0
2013-01-04     3.0
2013-01-05     5.0
2013-01-06    NaN
Freq: D, dtype: float64

In [65]: df.sub(s, axis="index")
Out[65]:
```

	A	B	C	D	F
2013-01-01	NaN	NaN	NaN	NaN	NaN
2013-01-02	NaN	NaN	NaN	NaN	NaN
2013-01-03	-1.139201	-3.090971	-0.024671	4.0	1.0
2013-01-04	-2.751724	-4.001166	-2.873309	2.0	0.0
2013-01-05	-4.720643	-5.002987	-5.566766	0.0	-1.0
2013-01-06	NaN	NaN	NaN	NaN	NaN

Apply

Applying functions to the data:

```
In [66]: df.apply(np.cumsum)
Out[66]:
```

	A	B	C	D	F
2013-01-01	0.000000	0.000000	-0.945427	5	NaN
2013-01-02	-0.177498	1.487290	-1.403760	10	1.0
2013-01-03	-0.316699	-0.603681	-0.428431	15	3.0
2013-01-04	-0.068423	-1.604847	-0.301740	20	6.0
2013-01-05	0.210934	-1.607834	-0.868506	25	10.0
2013-01-06	-0.950702	-2.202828	-0.145855	30	15.0

```
In [67]: df.apply(lambda x: x.max() - x.min())
Out[67]:
```

	A	B	C	D	F
A	1.440992				
B	3.578262				
C	1.920756				
D	0.000000				
F	4.000000				

```
dtype: float64
```

Histogramming

See more at [Histogramming and Discretization](#).

```
In [68]: s = pd.Series(np.random.randint(0, 7, size=10))
In [69]: s
Out[69]:
```

0	1
1	2
2	0
3	5
4	1
5	0
6	0
7	3
8	5
9	0

```
dtype: int64

In [70]: s.value_counts()
Out[70]:
```

0	4
1	2
5	2
2	1
3	1

```
dtype: int64
```

String Methods

Series is equipped with a set of string processing methods in the `str` attribute that make it easy to operate on each element of the array, as in the code snippet below. Note that pattern-matching in `str` generally uses [regular expressions](#) by default (and in some cases always uses them). See more at [Vectorized String Methods](#).

```
In [71]: s = pd.Series(["A", "B", "C", "Aaba", "Baca", np.nan, "CABA", "dog", "cat"])

In [72]: s.str.lower()
Out[72]:
0      a
1      b
2      c
3    aaba
4    baca
5     NaN
6    caba
7     dog
8     cat
dtype: object
```

2.1.6 Merge

Concat

pandas provides various facilities for easily combining together Series and DataFrame objects with various kinds of set logic for the indexes and relational algebra functionality in the case of join / merge-type operations.

See the [Merging section](#).

Concatenating pandas objects together with `concat()`:

```
In [73]: df = pd.DataFrame(np.random.randn(10, 4))

In [74]: df
Out[74]:
   0         1         2         3
0  1.689662 -1.311349 -0.229433  0.687357
1 -1.225525 -1.555269  1.647132  2.163551
2  1.133407  0.362450  0.238007  0.166870
3 -0.050427  0.091661 -0.532976  0.071773
4  1.633323  1.360126 -0.509585  0.008281
5 -0.874520 -1.203754  0.542787 -0.488790
6  1.960843 -0.555579  0.762490  0.614574
7  0.580310 -0.200614  0.501313 -1.571428
8 -0.818894 -1.714748 -0.110789 -1.783606
9  0.257285 -0.300526  1.599902  0.247784

# break it into pieces
In [75]: pieces = [df[:3], df[3:7], df[7:]]

In [76]: pd.concat(pieces)
Out[76]:
   0         1         2         3
0  1.689662 -1.311349 -0.229433  0.687357
1 -1.225525 -1.555269  1.647132  2.163551
```

(continues on next page)

(continued from previous page)

```

2  1.133407  0.362450  0.238007  0.166870
3 -0.050427  0.091661 -0.532976  0.071773
4  1.633323  1.360126 -0.509585  0.008281
5 -0.874520 -1.203754  0.542787 -0.488790
6  1.960843 -0.555579  0.762490  0.614574
7  0.580310 -0.200614  0.501313 -1.571428
8 -0.818894 -1.714748 -0.110789 -1.783606
9  0.257285 -0.300526  1.599902  0.247784

```

Note: Adding a column to a DataFrame is relatively fast. However, adding a row requires a copy, and may be expensive. We recommend passing a pre-built list of records to the DataFrame constructor instead of building a DataFrame by iteratively appending records to it. See [Appending to dataframe](#) for more.

Join

SQL style merges. See the [Database style joining](#) section.

```

In [77]: left = pd.DataFrame({"key": ["foo", "foo"], "lval": [1, 2]})

In [78]: right = pd.DataFrame({"key": ["foo", "foo"], "rval": [4, 5]})

In [79]: left
Out[79]:
   key  lval
0  foo     1
1  foo     2

In [80]: right
Out[80]:
   key  rval
0  foo     4
1  foo     5

In [81]: pd.merge(left, right, on="key")
Out[81]:
   key  lval  rval
0  foo     1     4
1  foo     1     5
2  foo     2     4
3  foo     2     5

```

Another example that can be given is:

```

In [82]: left = pd.DataFrame({"key": ["foo", "bar"], "lval": [1, 2]})

In [83]: right = pd.DataFrame({"key": ["foo", "bar"], "rval": [4, 5]})

In [84]: left
Out[84]:
   key  lval
0  foo     1
1  bar     2

```

(continues on next page)

(continued from previous page)

```
In [85]: right
Out[85]:
   key  rval
0  foo     4
1  bar     5

In [86]: pd.merge(left, right, on="key")
Out[86]:
   key  lval  rval
0  foo     1     4
1  bar     2     5
```

2.1.7 Grouping

By “group by” we are referring to a process involving one or more of the following steps:

- **Splitting** the data into groups based on some criteria
- **Applying** a function to each group independently
- **Combining** the results into a data structure

See the *Grouping section*.

```
In [87]: df = pd.DataFrame(
...:     {
...:         "A": ["foo", "bar", "foo", "bar", "foo", "bar", "foo", "foo"],
...:         "B": ["one", "one", "two", "three", "two", "two", "one", "three"],
...:         "C": np.random.randn(8),
...:         "D": np.random.randn(8),
...:     }
...: )
...:

In [88]: df
Out[88]:
   A      B      C      D
0  foo  one  1.355302 -0.578010
1  bar  one  1.433322  0.238339
2  foo  two  0.158418  1.066454
3  bar three  0.196638  0.327690
4  foo  two -1.244050  2.601324
5  bar  two  0.986477 -0.845308
6  foo  one -0.770289  0.183200
7  foo three  0.443799  0.418662
```

Grouping and then applying the `sum()` function to the resulting groups.

```
In [89]: df.groupby("A").sum()
Out[89]:
      C      D
A
bar  2.616436 -0.279278
foo -0.056820  3.691630
```

Grouping by multiple columns forms a hierarchical index, and again we can apply the `sum()` function.

```
In [90]: df.groupby(["A", "B"]).sum()
Out[90]:
```

		C	D
A	B		
bar	one	1.433322	0.238339
	three	0.196638	0.327690
	two	0.986477	-0.845308
foo	one	0.585014	-0.394810
	three	0.443799	0.418662
	two	-1.085632	3.667778

2.1.8 Reshaping

See the sections on *Hierarchical Indexing* and *Reshaping*.

Stack

```
In [91]: tuples = list(
....:     zip(
....:         *[
....:             ["bar", "bar", "baz", "baz", "foo", "foo", "qux", "qux"],
....:             ["one", "two", "one", "two", "one", "two", "one", "two"],
....:         ]
....:     )
....: )
....:

In [92]: index = pd.MultiIndex.from_tuples(tuples, names=["first", "second"])

In [93]: df = pd.DataFrame(np.random.randn(8, 2), index=index, columns=["A", "B"])

In [94]: df2 = df[:4]

In [95]: df2
Out[95]:
```

		A	B
first	second		
bar	one	-0.908885	0.134150
	two	-0.322778	-0.084722
baz	one	-0.095210	-2.256921
	two	2.615210	0.214028

The `stack()` method “compresses” a level in the `DataFrame`’s columns.

```
In [96]: stacked = df2.stack()

In [97]: stacked
Out[97]:
```

first	second		
bar	one	A	-0.908885
		B	0.134150
	two	A	-0.322778
		B	-0.084722
baz	one	A	-0.095210

(continues on next page)

(continued from previous page)

```

            B    -2.256921
two        A     2.615210
            B     0.214028
dtype: float64

```

With a “stacked” DataFrame or Series (having a MultiIndex as the index), the inverse operation of `stack()` is `unstack()`, which by default unstacks the **last level**:

```
In [98]: stacked.unstack()
```

```
Out [98]:
```

```

              A          B
first second
bar  one  -0.908885  0.134150
     two  -0.322778 -0.084722
baz   one  -0.095210 -2.256921
     two   2.615210  0.214028

```

```
In [99]: stacked.unstack(1)
```

```
Out [99]:
```

```

second      one      two
first
bar  A -0.908885 -0.322778
     B  0.134150 -0.084722
baz  A -0.095210  2.615210
     B -2.256921  0.214028

```

```
In [100]: stacked.unstack(0)
```

```
Out [100]:
```

```

first      bar      baz
second
one    A -0.908885 -0.095210
       B  0.134150 -2.256921
two    A -0.322778  2.615210
       B -0.084722  0.214028

```

Pivot tables

See the section on *Pivot Tables*.

```

In [101]: df = pd.DataFrame(
.....:     {
.....:         "A": ["one", "one", "two", "three"] * 3,
.....:         "B": ["A", "B", "C"] * 4,
.....:         "C": ["foo", "foo", "foo", "bar", "bar", "bar"] * 2,
.....:         "D": np.random.randn(12),
.....:         "E": np.random.randn(12),
.....:     }
.....: )
.....:

```

```
In [102]: df
```

```
Out [102]:
```

```

      A  B  C      D      E
0  one  A  foo  0.745325 -1.492997
1  one  B  foo -1.297547 -0.339886

```

(continues on next page)

(continued from previous page)

```

2    two  C  foo -0.964388  0.482778
3   three A  bar  1.079703  1.604137
4    one  B  bar  0.675036  0.560282
5    one  C  bar  0.408994  0.558150
6    two  A  foo  0.261266 -0.149187
7   three B  foo  1.012788 -1.226392
8    one  C  foo -0.105359  1.401395
9    one  A  bar  1.180264 -0.625340
10   two  B  bar  1.648181 -1.047040
11  three  C  bar  0.792630 -1.237315

```

We can produce pivot tables from this data very easily:

```

In [103]: pd.pivot_table(df, values="D", index=["A", "B"], columns=["C"])
Out[103]:
C          bar          foo
A    B
one  A  1.180264  0.745325
     B  0.675036 -1.297547
     C  0.408994 -0.105359
three A  1.079703         NaN
     B         NaN  1.012788
     C  0.792630         NaN
two   A         NaN  0.261266
     B  1.648181         NaN
     C         NaN -0.964388

```

2.1.9 Time series

pandas has simple, powerful, and efficient functionality for performing resampling operations during frequency conversion (e.g., converting secondly data into 5-minutely data). This is extremely common in, but not limited to, financial applications. See the [Time Series section](#).

```

In [104]: rng = pd.date_range("1/1/2012", periods=100, freq="S")

In [105]: ts = pd.Series(np.random.randint(0, 500, len(rng)), index=rng)

In [106]: ts.resample("5Min").sum()
Out[106]:
2012-01-01    26891
Freq: 5T, dtype: int64

```

Time zone representation:

```

In [107]: rng = pd.date_range("3/6/2012 00:00", periods=5, freq="D")

In [108]: ts = pd.Series(np.random.randn(len(rng)), rng)

In [109]: ts
Out[109]:
2012-03-06    0.174529
2012-03-07    0.457967
2012-03-08   -0.314052
2012-03-09   -0.072789
2012-03-10   -1.091168

```

(continues on next page)

(continued from previous page)

```

Freq: D, dtype: float64

In [110]: ts_utc = ts.tz_localize("UTC")

In [111]: ts_utc
Out[111]:
2012-03-06 00:00:00+00:00    0.174529
2012-03-07 00:00:00+00:00    0.457967
2012-03-08 00:00:00+00:00   -0.314052
2012-03-09 00:00:00+00:00   -0.072789
2012-03-10 00:00:00+00:00   -1.091168
Freq: D, dtype: float64

```

Converting to another time zone:

```

In [112]: ts_utc.tz_convert("US/Eastern")
Out[112]:
2012-03-05 19:00:00-05:00    0.174529
2012-03-06 19:00:00-05:00    0.457967
2012-03-07 19:00:00-05:00   -0.314052
2012-03-08 19:00:00-05:00   -0.072789
2012-03-09 19:00:00-05:00   -1.091168
Freq: D, dtype: float64

```

Converting between time span representations:

```

In [113]: rng = pd.date_range("1/1/2012", periods=5, freq="M")

In [114]: ts = pd.Series(np.random.randn(len(rng)), index=rng)

In [115]: ts
Out[115]:
2012-01-31    1.262722
2012-02-29    0.203150
2012-03-31    0.129981
2012-04-30    0.376551
2012-05-31    0.103878
Freq: M, dtype: float64

In [116]: ps = ts.to_period()

In [117]: ps
Out[117]:
2012-01    1.262722
2012-02    0.203150
2012-03    0.129981
2012-04    0.376551
2012-05    0.103878
Freq: M, dtype: float64

In [118]: ps.to_timestamp()
Out[118]:
2012-01-01    1.262722
2012-02-01    0.203150
2012-03-01    0.129981
2012-04-01    0.376551
2012-05-01    0.103878

```

(continues on next page)

(continued from previous page)

```
Freq: MS, dtype: float64
```

Converting between period and timestamp enables some convenient arithmetic functions to be used. In the following example, we convert a quarterly frequency with year ending in November to 9am of the end of the month following the quarter end:

```
In [119]: prng = pd.period_range("1990Q1", "2000Q4", freq="Q-NOV")
In [120]: ts = pd.Series(np.random.randn(len(prng)), prng)
In [121]: ts.index = (prng.asfreq("M", "e") + 1).asfreq("H", "s") + 9
In [122]: ts.head()
Out[122]:
1990-03-01 09:00    0.124480
1990-06-01 09:00    0.406012
1990-09-01 09:00    0.203973
1990-12-01 09:00    0.355695
1991-03-01 09:00    2.962668
Freq: H, dtype: float64
```

2.1.10 Categoricals

pandas can include categorical data in a DataFrame. For full docs, see the *categorical introduction* and the *API documentation*.

```
In [123]: df = pd.DataFrame(
.....:     {"id": [1, 2, 3, 4, 5, 6], "raw_grade": ["a", "b", "b", "a", "a", "e"]}
.....: )
.....:
```

Convert the raw grades to a categorical data type.

```
In [124]: df["grade"] = df["raw_grade"].astype("category")
In [125]: df["grade"]
Out[125]:
0    a
1    b
2    b
3    a
4    a
5    e
Name: grade, dtype: category
Categories (3, object): ['a', 'b', 'e']
```

Rename the categories to more meaningful names (assigning to `Series.cat.categories()` is in place!).

```
In [126]: df["grade"].cat.categories = ["very good", "good", "very bad"]
```

Reorder the categories and simultaneously add the missing categories (methods under `Series.cat()` return a new Series by default).

```
In [127]: df["grade"] = df["grade"].cat.set_categories(
.....:     ["very bad", "bad", "medium", "good", "very good"]
.....: )
.....:

In [128]: df["grade"]
Out[128]:
0    very good
1         good
2         good
3    very good
4    very good
5    very bad
Name: grade, dtype: category
Categories (5, object): ['very bad', 'bad', 'medium', 'good', 'very good']
```

Sorting is per order in the categories, not lexical order.

```
In [129]: df.sort_values(by="grade")
Out[129]:
   id raw_grade  grade
5   6         e  very bad
1   2         b    good
2   3         b    good
0   1         a  very good
3   4         a  very good
4   5         a  very good
```

Grouping by a categorical column also shows empty categories.

```
In [130]: df.groupby("grade").size()
Out[130]:
grade
very bad    1
bad         0
medium      0
good        2
very good   3
dtype: int64
```

2.1.11 Plotting

See the *Plotting* docs.

We use the standard convention for referencing the matplotlib API:

```
In [131]: import matplotlib.pyplot as plt

In [132]: plt.close("all")
```

The `close()` method is used to *close* a figure window.


```
In [133]: ts = pd.Series(np.random.randn(1000), index=pd.date_range("1/1/2000",
↳ periods=1000))

In [134]: ts = ts.cumsum()
```

(continues on next page)

(continued from previous page)

```
In [135]: ts.plot();
```



On a `DataFrame`, the `plot()` method is a convenience to plot all of the columns with labels:

```
In [136]: df = pd.DataFrame(
.....:     np.random.randn(1000, 4), index=ts.index, columns=["A", "B", "C", "D"]
.....: )
.....:

In [137]: df = df.cumsum()

In [138]: plt.figure();

In [139]: df.plot();

In [140]: plt.legend(loc='best');
```



2.1.12 Getting data in/out

CSV

Writing to a csv file.

```
In [141]: df.to_csv("foo.csv")
```

Reading from a csv file.

```
In [142]: pd.read_csv("foo.csv")
```

Out[142]:

	Unnamed: 0	A	B	C	D
0	2000-01-01	-1.065194	-0.002634	-0.760958	-0.031015
1	2000-01-02	-1.215840	-0.589713	0.060526	-1.667720
2	2000-01-03	-1.018093	-0.473960	-0.930696	-0.735057
3	2000-01-04	-0.574559	-2.144053	-0.009284	-1.372000
4	2000-01-05	-1.215664	-3.426502	-0.283470	-0.781581
...
995	2002-09-22	2.281268	-24.942834	52.881026	27.864387
996	2002-09-23	3.271586	-25.498553	52.312461	28.249057
997	2002-09-24	1.913191	-26.822220	52.965859	27.642595
998	2002-09-25	1.690795	-26.325976	52.159921	26.247726

(continues on next page)

(continued from previous page)

```
999 2002-09-26 -0.098981 -26.778316 52.012625 25.801426

[1000 rows x 5 columns]
```

HDF5

Reading and writing to *HDFStores*.

Writing to a HDF5 Store.

```
In [143]: df.to_hdf("foo.h5", "df")
```

Reading from a HDF5 Store.

```
In [144]: pd.read_hdf("foo.h5", "df")
Out[144]:
```

	A	B	C	D
2000-01-01	-1.065194	-0.002634	-0.760958	-0.031015
2000-01-02	-1.215840	-0.589713	0.060526	-1.667720
2000-01-03	-1.018093	-0.473960	-0.930696	-0.735057
2000-01-04	-0.574559	-2.144053	-0.009284	-1.372000
2000-01-05	-1.215664	-3.426502	-0.283470	-0.781581
...
2002-09-22	2.281268	-24.942834	52.881026	27.864387
2002-09-23	3.271586	-25.498553	52.312461	28.249057
2002-09-24	1.913191	-26.822220	52.965859	27.642595
2002-09-25	1.690795	-26.325976	52.159921	26.247726
2002-09-26	-0.098981	-26.778316	52.012625	25.801426

```
[1000 rows x 4 columns]
```

Excel

Reading and writing to *MS Excel*.

Writing to an excel file.

```
In [145]: df.to_excel("foo.xlsx", sheet_name="Sheet1")
```

Reading from an excel file.

```
In [146]: pd.read_excel("foo.xlsx", "Sheet1", index_col=None, na_values=["NA"])
Out[146]:
```

	Unnamed: 0	A	B	C	D
0	2000-01-01	-1.065194	-0.002634	-0.760958	-0.031015
1	2000-01-02	-1.215840	-0.589713	0.060526	-1.667720
2	2000-01-03	-1.018093	-0.473960	-0.930696	-0.735057
3	2000-01-04	-0.574559	-2.144053	-0.009284	-1.372000
4	2000-01-05	-1.215664	-3.426502	-0.283470	-0.781581
...
995	2002-09-22	2.281268	-24.942834	52.881026	27.864387
996	2002-09-23	3.271586	-25.498553	52.312461	28.249057
997	2002-09-24	1.913191	-26.822220	52.965859	27.642595
998	2002-09-25	1.690795	-26.325976	52.159921	26.247726

(continues on next page)

(continued from previous page)

```
999 2002-09-26 -0.098981 -26.778316 52.012625 25.801426
[1000 rows x 5 columns]
```

2.1.13 Gotchas

If you are attempting to perform an operation you might see an exception like:

```
>>> if pd.Series([False, True, False]):
...     print("I was true")
Traceback
...
ValueError: The truth value of an array is ambiguous. Use a.empty, a.any() or a.all().
```

See [Comparisons](#) for an explanation and what to do.

See [Gotchas](#) as well.

2.2 Intro to data structures

We'll start with a quick, non-comprehensive overview of the fundamental data structures in pandas to get you started. The fundamental behavior about data types, indexing, and axis labeling / alignment apply across all of the objects. To get started, import NumPy and load pandas into your namespace:

```
In [1]: import numpy as np
In [2]: import pandas as pd
```

Here is a basic tenet to keep in mind: **data alignment is intrinsic**. The link between labels and data will not be broken unless done so explicitly by you.

We'll give a brief intro to the data structures, then consider all of the broad categories of functionality and methods in separate sections.

2.2.1 Series

Series is a one-dimensional labeled array capable of holding any data type (integers, strings, floating point numbers, Python objects, etc.). The axis labels are collectively referred to as the **index**. The basic method to create a Series is to call:

```
>>> s = pd.Series(data, index=index)
```

Here, data can be many different things:

- a Python dict
- an ndarray
- a scalar value (like 5)

The passed **index** is a list of axis labels. Thus, this separates into a few cases depending on what **data** is:

From ndarray

If data is an ndarray, **index** must be the same length as **data**. If no index is passed, one will be created having values `[0, ..., len(data) - 1]`.

```
In [3]: s = pd.Series(np.random.randn(5), index=["a", "b", "c", "d", "e"])

In [4]: s
Out[4]:
a    0.469112
b   -0.282863
c   -1.509059
d   -1.135632
e    1.212112
dtype: float64

In [5]: s.index
Out[5]: Index(['a', 'b', 'c', 'd', 'e'], dtype='object')

In [6]: pd.Series(np.random.randn(5))
Out[6]:
0   -0.173215
1    0.119209
2   -1.044236
3   -0.861849
4   -2.104569
dtype: float64
```

Note: pandas supports non-unique index values. If an operation that does not support duplicate index values is attempted, an exception will be raised at that time. The reason for being lazy is nearly all performance-based (there are many instances in computations, like parts of GroupBy, where the index is not used).

From dict

Series can be instantiated from dicts:

```
In [7]: d = {"b": 1, "a": 0, "c": 2}

In [8]: pd.Series(d)
Out[8]:
b    1
a    0
c    2
dtype: int64
```

Note: When the data is a dict, and an index is not passed, the `Series` index will be ordered by the dict's insertion order, if you're using Python version `>= 3.6` and pandas version `>= 0.23`.

If you're using Python `< 3.6` or pandas `< 0.23`, and an index is not passed, the `Series` index will be the lexically ordered list of dict keys.

In the example above, if you were on a Python version lower than 3.6 or a pandas version lower than 0.23, the `Series` would be ordered by the lexical order of the dict keys (i.e. `['a', 'b', 'c']` rather than `['b', 'a', 'c']`).

If an index is passed, the values in data corresponding to the labels in the index will be pulled out.

```
In [9]: d = {"a": 0.0, "b": 1.0, "c": 2.0}

In [10]: pd.Series(d)
Out[10]:
a    0.0
b    1.0
c    2.0
dtype: float64

In [11]: pd.Series(d, index=["b", "c", "d", "a"])
Out[11]:
b    1.0
c    2.0
d    NaN
a    0.0
dtype: float64
```

Note: NaN (not a number) is the standard missing data marker used in pandas.

From scalar value

If data is a scalar value, an index must be provided. The value will be repeated to match the length of **index**.

```
In [12]: pd.Series(5.0, index=["a", "b", "c", "d", "e"])
Out[12]:
a    5.0
b    5.0
c    5.0
d    5.0
e    5.0
dtype: float64
```

Series is ndarray-like

Series acts very similarly to a ndarray, and is a valid argument to most NumPy functions. However, operations such as slicing will also slice the index.

```
In [13]: s[0]
Out[13]: 0.4691122999071863

In [14]: s[:3]
Out[14]:
a    0.469112
b   -0.282863
c   -1.509059
dtype: float64

In [15]: s[s > s.median()]
Out[15]:
a    0.469112
e    1.212112
dtype: float64

In [16]: s[[4, 3, 1]]
```

(continues on next page)

(continued from previous page)

```
Out [16]:  
e    1.212112  
d   -1.135632  
b   -0.282863  
dtype: float64  
  
In [17]: np.exp(s)  
Out [17]:  
a    1.598575  
b    0.753623  
c    0.221118  
d    0.321219  
e    3.360575  
dtype: float64
```

Note: We will address array-based indexing like `s[[4, 3, 1]]` in [section on indexing](#).

Like a NumPy array, a pandas Series has a *dtype*.

```
In [18]: s.dtype  
Out [18]: dtype('float64')
```

This is often a NumPy dtype. However, pandas and 3rd-party libraries extend NumPy's type system in a few places, in which case the dtype would be an *ExtensionDtype*. Some examples within pandas are [Categorical data](#) and [Nullable integer data type](#). See [dtypes](#) for more.

If you need the actual array backing a Series, use `Series.array`.

```
In [19]: s.array  
Out [19]:  
<PandasArray>  
[ 0.4691122999071863, -0.2828633443286633, -1.5090585031735124,  
 -1.1356323710171934,  1.2121120250208506]  
Length: 5, dtype: float64
```

Accessing the array can be useful when you need to do some operation without the index (to disable *automatic alignment*, for example).

`Series.array` will always be an *ExtensionArray*. Briefly, an *ExtensionArray* is a thin wrapper around one or more *concrete* arrays like a `numpy.ndarray`. pandas knows how to take an *ExtensionArray* and store it in a Series or a column of a DataFrame. See [dtypes](#) for more.

While Series is ndarray-like, if you need an *actual* ndarray, then use `Series.to_numpy()`.

```
In [20]: s.to_numpy()  
Out [20]: array([ 0.4691, -0.2829, -1.5091, -1.1356,  1.2121])
```

Even if the Series is backed by a *ExtensionArray*, `Series.to_numpy()` will return a NumPy ndarray.

Series is dict-like

A Series is like a fixed-size dict in that you can get and set values by index label:

```
In [21]: s["a"]
Out[21]: 0.4691122999071863

In [22]: s["e"] = 12.0

In [23]: s
Out[23]:
a    0.469112
b   -0.282863
c   -1.509059
d   -1.135632
e   12.000000
dtype: float64

In [24]: "e" in s
Out[24]: True

In [25]: "f" in s
Out[25]: False
```

If a label is not contained, an exception is raised:

```
>>> s["f"]
KeyError: 'f'
```

Using the `get` method, a missing label will return `None` or specified default:

```
In [26]: s.get("f")

In [27]: s.get("f", np.nan)
Out[27]: nan
```

See also the [section on attribute access](#).

Vectorized operations and label alignment with Series

When working with raw NumPy arrays, looping through value-by-value is usually not necessary. The same is true when working with Series in pandas. Series can also be passed into most NumPy methods expecting an ndarray.

```
In [28]: s + s
Out[28]:
a    0.938225
b   -0.565727
c   -3.018117
d   -2.271265
e   24.000000
dtype: float64

In [29]: s * 2
Out[29]:
a    0.938225
b   -0.565727
```

(continues on next page)

(continued from previous page)

```
c    -3.018117
d    -2.271265
e     24.000000
dtype: float64

In [30]: np.exp(s)
Out[30]:
a         1.598575
b         0.753623
c         0.221118
d         0.321219
e    162754.791419
dtype: float64
```

A key difference between Series and ndarray is that operations between Series automatically align the data based on label. Thus, you can write computations without giving consideration to whether the Series involved have the same labels.

```
In [31]: s[1:] + s[:-1]
Out[31]:
a         NaN
b    -0.565727
c    -3.018117
d    -2.271265
e         NaN
dtype: float64
```

The result of an operation between unaligned Series will have the **union** of the indexes involved. If a label is not found in one Series or the other, the result will be marked as missing NaN. Being able to write code without doing any explicit data alignment grants immense freedom and flexibility in interactive data analysis and research. The integrated data alignment features of the pandas data structures set pandas apart from the majority of related tools for working with labeled data.

Note: In general, we chose to make the default result of operations between differently indexed objects yield the **union** of the indexes in order to avoid loss of information. Having an index label, though the data is missing, is typically important information as part of a computation. You of course have the option of dropping labels with missing data via the **dropna** function.

Name attribute

Series can also have a name attribute:

```
In [32]: s = pd.Series(np.random.randn(5), name="something")

In [33]: s
Out[33]:
0    -0.494929
1     1.071804
2     0.721555
3    -0.706771
4    -1.039575
Name: something, dtype: float64
```

(continues on next page)

(continued from previous page)

```
In [34]: s.name
Out[34]: 'something'
```

The Series `name` will be assigned automatically in many cases, in particular when taking 1D slices of `DataFrame` as you will see below.

You can rename a Series with the `pandas.Series.rename()` method.

```
In [35]: s2 = s.rename("different")

In [36]: s2.name
Out[36]: 'different'
```

Note that `s` and `s2` refer to different objects.

2.2.2 DataFrame

DataFrame is a 2-dimensional labeled data structure with columns of potentially different types. You can think of it like a spreadsheet or SQL table, or a dict of Series objects. It is generally the most commonly used pandas object. Like Series, `DataFrame` accepts many different kinds of input:

- Dict of 1D ndarrays, lists, dicts, or Series
- 2-D `numpy.ndarray`
- `Structured or record ndarray`
- A Series
- Another `DataFrame`

Along with the data, you can optionally pass **index** (row labels) and **columns** (column labels) arguments. If you pass an index and / or columns, you are guaranteeing the index and / or columns of the resulting `DataFrame`. Thus, a dict of Series plus a specific index will discard all data not matching up to the passed index.

If axis labels are not passed, they will be constructed from the input data based on common sense rules.

Note: When the data is a dict, and `columns` is not specified, the `DataFrame` columns will be ordered by the dict's insertion order, if you are using Python version `>= 3.6` and pandas `>= 0.23`.

If you are using Python `< 3.6` or pandas `< 0.23`, and `columns` is not specified, the `DataFrame` columns will be the lexically ordered list of dict keys.

From dict of Series or dicts

The resulting **index** will be the **union** of the indexes of the various Series. If there are any nested dicts, these will first be converted to Series. If no columns are passed, the columns will be the ordered list of dict keys.

```
In [37]: d = {
.....:     "one": pd.Series([1.0, 2.0, 3.0], index=["a", "b", "c"]),
.....:     "two": pd.Series([1.0, 2.0, 3.0, 4.0], index=["a", "b", "c", "d"]),
.....: }
.....:

In [38]: df = pd.DataFrame(d)
```

(continues on next page)

(continued from previous page)

```

In [39]: df
Out[39]:
   one  two
a  1.0  1.0
b  2.0  2.0
c  3.0  3.0
d  NaN  4.0

In [40]: pd.DataFrame(d, index=["d", "b", "a"])
Out[40]:
   one  two
d  NaN  4.0
b  2.0  2.0
a  1.0  1.0

In [41]: pd.DataFrame(d, index=["d", "b", "a"], columns=["two", "three"])
Out[41]:
   two three
d  4.0   NaN
b  2.0   NaN
a  1.0   NaN

```

The row and column labels can be accessed respectively by accessing the **index** and **columns** attributes:

Note: When a particular set of columns is passed along with a dict of data, the passed columns override the keys in the dict.

```

In [42]: df.index
Out[42]: Index(['a', 'b', 'c', 'd'], dtype='object')

In [43]: df.columns
Out[43]: Index(['one', 'two'], dtype='object')

```

From dict of ndarrays / lists

The ndarrays must all be the same length. If an index is passed, it must clearly also be the same length as the arrays. If no index is passed, the result will be `range(n)`, where `n` is the array length.

```

In [44]: d = {"one": [1.0, 2.0, 3.0, 4.0], "two": [4.0, 3.0, 2.0, 1.0]}

In [45]: pd.DataFrame(d)
Out[45]:
   one  two
0  1.0  4.0
1  2.0  3.0
2  3.0  2.0
3  4.0  1.0

In [46]: pd.DataFrame(d, index=["a", "b", "c", "d"])
Out[46]:
   one  two
a  1.0  4.0

```

(continues on next page)

(continued from previous page)

```
b  2.0  3.0
c  3.0  2.0
d  4.0  1.0
```

From structured or record array

This case is handled identically to a dict of arrays.

```
In [47]: data = np.zeros((2,), dtype=[("A", "i4"), ("B", "f4"), ("C", "a10")])

In [48]: data[:] = [(1, 2.0, "Hello"), (2, 3.0, "World")]

In [49]: pd.DataFrame(data)
Out[49]:
   A    B      C
0  1  2.0 b'Hello'
1  2  3.0 b'World'

In [50]: pd.DataFrame(data, index=["first", "second"])
Out[50]:
      A    B      C
first  1  2.0 b'Hello'
second 2  3.0 b'World'

In [51]: pd.DataFrame(data, columns=["C", "A", "B"])
Out[51]:
      C    A    B
0 b'Hello'  1  2.0
1 b'World'  2  3.0
```

Note: DataFrame is not intended to work exactly like a 2-dimensional NumPy ndarray.

From a list of dicts

```
In [52]: data2 = [{"a": 1, "b": 2}, {"a": 5, "b": 10, "c": 20}]

In [53]: pd.DataFrame(data2)
Out[53]:
   a  b      c
0  1  2   NaN
1  5 10  20.0

In [54]: pd.DataFrame(data2, index=["first", "second"])
Out[54]:
      a  b      c
first  1  2   NaN
second 5 10  20.0

In [55]: pd.DataFrame(data2, columns=["a", "b"])
Out[55]:
   a  b
```

(continues on next page)

(continued from previous page)

```
0  1   2
1  5  10
```

From a dict of tuples

You can automatically create a MultiIndexed frame by passing a tuples dictionary.

```
In [56]: pd.DataFrame(
...:     {
...:         ("a", "b"): {"A", "B": 1, ("A", "C"): 2},
...:         ("a", "a"): {"A", "C": 3, ("A", "B"): 4},
...:         ("a", "c"): {"A", "B": 5, ("A", "C"): 6},
...:         ("b", "a"): {"A", "C": 7, ("A", "B"): 8},
...:         ("b", "b"): {"A", "D": 9, ("A", "B"): 10},
...:     }
...: )
Out[56]:
```

		a		b		
		b	a	c	a	b
A	B	1.0	4.0	5.0	8.0	10.0
	C	2.0	3.0	6.0	7.0	NaN
	D	NaN	NaN	NaN	NaN	9.0

From a Series

The result will be a DataFrame with the same index as the input Series, and with one column whose name is the original name of the Series (only if no other column name provided).

From a list of namedtuples

The field names of the first `namedtuple` in the list determine the columns of the DataFrame. The remaining `namedtuples` (or tuples) are simply unpacked and their values are fed into the rows of the DataFrame. If any of those tuples is shorter than the first `namedtuple` then the later columns in the corresponding row are marked as missing values. If any are longer than the first `namedtuple`, a `ValueError` is raised.

```
In [57]: from collections import namedtuple

In [58]: Point = namedtuple("Point", "x y")

In [59]: pd.DataFrame([Point(0, 0), Point(0, 3), (2, 3)])
Out[59]:
```

	x	y
0	0	0
1	0	3
2	2	3

```
In [60]: Point3D = namedtuple("Point3D", "x y z")

In [61]: pd.DataFrame([Point3D(0, 0, 0), Point3D(0, 3, 5), Point(2, 3)])
Out[61]:
```

	x	y	z
0	0	0	0
1	0	3	5
2	2	3	NaN

(continues on next page)

(continued from previous page)

```
0  0  0  0.0
1  0  3  5.0
2  2  3  NaN
```

From a list of dataclasses

New in version 1.1.0.

Data Classes as introduced in [PEP557](#), can be passed into the DataFrame constructor. Passing a list of dataclasses is equivalent to passing a list of dictionaries.

Please be aware, that all values in the list should be dataclasses, mixing types in the list would result in a `TypeError`.

```
In [62]: from dataclasses import make_dataclass

In [63]: Point = make_dataclass("Point", [("x", int), ("y", int)])

In [64]: pd.DataFrame([Point(0, 0), Point(0, 3), Point(2, 3)])
Out[64]:
   x  y
0  0  0
1  0  3
2  2  3
```

Missing data

Much more will be said on this topic in the [Missing data](#) section. To construct a DataFrame with missing data, we use `np.nan` to represent missing values. Alternatively, you may pass a `numpy.MaskedArray` as the data argument to the DataFrame constructor, and its masked entries will be considered missing.

Alternate constructors

DataFrame.from_dict

`DataFrame.from_dict` takes a dict of dicts or a dict of array-like sequences and returns a DataFrame. It operates like the DataFrame constructor except for the `orient` parameter which is `'columns'` by default, but which can be set to `'index'` in order to use the dict keys as row labels.

```
In [65]: pd.DataFrame.from_dict(dict([("A", [1, 2, 3]), ("B", [4, 5, 6])]))
Out[65]:
   A  B
0  1  4
1  2  5
2  3  6
```

If you pass `orient='index'`, the keys will be the row labels. In this case, you can also pass the desired column names:

```
In [66]: pd.DataFrame.from_dict(
....:     dict([("A", [1, 2, 3]), ("B", [4, 5, 6])]),
....:     orient="index",
....:     columns=["one", "two", "three"],
....: )
....:
Out[66]:
```

(continues on next page)

(continued from previous page)

	one	two	three
A	1	2	3
B	4	5	6

DataFrame.from_records

`DataFrame.from_records` takes a list of tuples or an ndarray with structured dtype. It works analogously to the normal `DataFrame` constructor, except that the resulting `DataFrame` index may be a specific field of the structured dtype. For example:

```
In [67]: data
Out[67]:
array([(1, 2., b'Hello'), (2, 3., b'World')],
      dtype=[('A', '<i4'), ('B', '<f4'), ('C', 'S10')])

In [68]: pd.DataFrame.from_records(data, index="C")
Out[68]:
```

	A	B
C		
b'Hello'	1	2.0
b'World'	2	3.0

Column selection, addition, deletion

You can treat a `DataFrame` semantically like a dict of like-indexed `Series` objects. Getting, setting, and deleting columns works with the same syntax as the analogous dict operations:

```
In [69]: df["one"]
Out[69]:
a    1.0
b    2.0
c    3.0
d    NaN
Name: one, dtype: float64

In [70]: df["three"] = df["one"] * df["two"]

In [71]: df["flag"] = df["one"] > 2

In [72]: df
Out[72]:
```

	one	two	three	flag
a	1.0	1.0	1.0	False
b	2.0	2.0	4.0	False
c	3.0	3.0	9.0	True
d	NaN	4.0	NaN	False

Columns can be deleted or popped like with a dict:

```
In [73]: del df["two"]

In [74]: three = df.pop("three")

In [75]: df
Out[75]:
```

(continues on next page)

(continued from previous page)

```

    one    flag
a  1.0  False
b  2.0  False
c  3.0   True
d  NaN  False

```

When inserting a scalar value, it will naturally be propagated to fill the column:

```
In [76]: df["foo"] = "bar"
```

```
In [77]: df
```

```
Out[77]:
```

```

    one    flag  foo
a  1.0  False  bar
b  2.0  False  bar
c  3.0   True  bar
d  NaN  False  bar

```

When inserting a Series that does not have the same index as the DataFrame, it will be conformed to the DataFrame's index:

```
In [78]: df["one_trunc"] = df["one"][:2]
```

```
In [79]: df
```

```
Out[79]:
```

```

    one    flag  foo  one_trunc
a  1.0  False  bar         1.0
b  2.0  False  bar         2.0
c  3.0   True  bar         NaN
d  NaN  False  bar         NaN

```

You can insert raw ndarrays but their length must match the length of the DataFrame's index.

By default, columns get inserted at the end. The `insert` function is available to insert at a particular location in the columns:

```
In [80]: df.insert(1, "bar", df["one"])
```

```
In [81]: df
```

```
Out[81]:
```

```

    one  bar    flag  foo  one_trunc
a  1.0  1.0  False  bar         1.0
b  2.0  2.0  False  bar         2.0
c  3.0  3.0   True  bar         NaN
d  NaN  NaN  False  bar         NaN

```

Assigning new columns in method chains

Inspired by `dplyr`'s `mutate` verb, `DataFrame` has an `assign()` method that allows you to easily create new columns that are potentially derived from existing columns.

```
In [82]: iris = pd.read_csv("data/iris.data")

In [83]: iris.head()
Out[83]:
```

	SepalLength	SepalWidth	PetalLength	PetalWidth	Name
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

```
In [84]: iris.assign(sepal_ratio=iris["SepalWidth"] / iris["SepalLength"]).head()
Out[84]:
```

	SepalLength	SepalWidth	PetalLength	PetalWidth	Name	sepal_ratio
0	5.1	3.5	1.4	0.2	Iris-setosa	0.686275
1	4.9	3.0	1.4	0.2	Iris-setosa	0.612245
2	4.7	3.2	1.3	0.2	Iris-setosa	0.680851
3	4.6	3.1	1.5	0.2	Iris-setosa	0.673913
4	5.0	3.6	1.4	0.2	Iris-setosa	0.720000

In the example above, we inserted a precomputed value. We can also pass in a function of one argument to be evaluated on the `DataFrame` being assigned to.


```
In [85]: iris.assign(sepal_ratio=lambda x: (x["SepalWidth"] / x["SepalLength"])).
↳ head()
Out[85]:
```

	SepalLength	SepalWidth	PetalLength	PetalWidth	Name	sepal_ratio
0	5.1	3.5	1.4	0.2	Iris-setosa	0.686275
1	4.9	3.0	1.4	0.2	Iris-setosa	0.612245
2	4.7	3.2	1.3	0.2	Iris-setosa	0.680851
3	4.6	3.1	1.5	0.2	Iris-setosa	0.673913
4	5.0	3.6	1.4	0.2	Iris-setosa	0.720000

`assign` **always** returns a copy of the data, leaving the original `DataFrame` untouched.

Passing a callable, as opposed to an actual value to be inserted, is useful when you don't have a reference to the `DataFrame` at hand. This is common when using `assign` in a chain of operations. For example, we can limit the `DataFrame` to just those observations with a Sepal Length greater than 5, calculate the ratio, and plot:

```
In [86]: (
.....:     iris.query("SepalLength > 5")
.....:     .assign(
.....:         SepalRatio=lambda x: x.SepalWidth / x.SepalLength,
.....:         PetalRatio=lambda x: x.PetalWidth / x.PetalLength,
.....:     )
.....:     .plot(kind="scatter", x="SepalRatio", y="PetalRatio")
.....: )
Out[86]: <AxesSubplot:xlabel='SepalRatio', ylabel='PetalRatio'>
```



Since a function is passed in, the function is computed on the DataFrame being assigned to. Importantly, this is the DataFrame that's been filtered to those rows with sepal length greater than 5. The filtering happens first, and then the ratio calculations. This is an example where we didn't have a reference to the *filtered* DataFrame available.

The function signature for `assign` is simply `**kwargs`. The keys are the column names for the new fields, and the values are either a value to be inserted (for example, a `Series` or NumPy array), or a function of one argument to be called on the DataFrame. A *copy* of the original DataFrame is returned, with the new values inserted.

Starting with Python 3.6 the order of `**kwargs` is preserved. This allows for *dependent* assignment, where an expression later in `**kwargs` can refer to a column created earlier in the same `assign()`.

```
In [87]: dfa = pd.DataFrame({"A": [1, 2, 3], "B": [4, 5, 6]})

In [88]: dfa.assign(C=lambda x: x["A"] + x["B"], D=lambda x: x["A"] + x["C"])
Out[88]:
```

	A	B	C	D
0	1	4	5	6
1	2	5	7	9
2	3	6	9	12

In the second expression, `x['C']` will refer to the newly created column, that's equal to `dfa['A'] + dfa['B']`.

Indexing / selection

The basics of indexing are as follows:

Operation	Syntax	Result
Select column	<code>df[col]</code>	Series
Select row by label	<code>df.loc[label]</code>	Series
Select row by integer location	<code>df.iloc[loc]</code>	Series
Slice rows	<code>df[5:10]</code>	DataFrame
Select rows by boolean vector	<code>df[bool_vec]</code>	DataFrame

Row selection, for example, returns a Series whose index is the columns of the DataFrame:

```
In [89]: df.loc["b"]
Out[89]:
one          2.0
bar          2.0
flag         False
foo          bar
one_trunc     2.0
Name: b, dtype: object

In [90]: df.iloc[2]
Out[90]:
one          3.0
bar          3.0
flag          True
foo          bar
one_trunc     NaN
Name: c, dtype: object
```

For a more exhaustive treatment of sophisticated label-based indexing and slicing, see the [section on indexing](#). We will address the fundamentals of reindexing / conforming to new sets of labels in the [section on reindexing](#).

Data alignment and arithmetic

Data alignment between DataFrame objects automatically align on **both the columns and the index (row labels)**. Again, the resulting object will have the union of the column and row labels.

```
In [91]: df = pd.DataFrame(np.random.randn(10, 4), columns=["A", "B", "C", "D"])

In [92]: df2 = pd.DataFrame(np.random.randn(7, 3), columns=["A", "B", "C"])

In [93]: df + df2
Out[93]:
      A          B          C  D
0  0.045691 -0.014138  1.380871 NaN
1 -0.955398 -1.501007  0.037181 NaN
2 -0.662690  1.534833 -0.859691 NaN
3 -2.452949  1.237274 -0.133712 NaN
4  1.414490  1.951676 -2.320422 NaN
5 -0.494922 -1.649727 -1.084601 NaN
6 -1.047551 -0.748572 -0.805479 NaN
7      NaN          NaN          NaN NaN
```

(continues on next page)

(continued from previous page)

8	NaN	NaN	NaN NaN
9	NaN	NaN	NaN NaN

When doing an operation between DataFrame and Series, the default behavior is to align the Series **index** on the DataFrame **columns**, thus **broadcasting** row-wise. For example:

```
In [94]: df - df.iloc[0]
Out[94]:
```

	A	B	C	D
0	0.000000	0.000000	0.000000	0.000000
1	-1.359261	-0.248717	-0.453372	-1.754659
2	0.253128	0.829678	0.010026	-1.991234
3	-1.311128	0.054325	-1.724913	-1.620544
4	0.573025	1.500742	-0.676070	1.367331
5	-1.741248	0.781993	-1.241620	-2.053136
6	-1.240774	-0.869551	-0.153282	0.000430
7	-0.743894	0.411013	-0.929563	-0.282386
8	-1.194921	1.320690	0.238224	-1.482644
9	2.293786	1.856228	0.773289	-1.446531

For explicit control over the matching and broadcasting behavior, see the section on *flexible binary operations*.

Operations with scalars are just as you would expect:

```
In [95]: df * 5 + 2
Out[95]:
```

	A	B	C	D
0	3.359299	-0.124862	4.835102	3.381160
1	-3.437003	-1.368449	2.568242	-5.392133
2	4.624938	4.023526	4.885230	-6.575010
3	-3.196342	0.146766	-3.789461	-4.721559
4	6.224426	7.378849	1.454750	10.217815
5	-5.346940	3.785103	-1.373001	-6.884519
6	-2.844569	-4.472618	4.068691	3.383309
7	-0.360173	1.930201	0.187285	1.969232
8	-2.615303	6.478587	6.026220	-4.032059
9	14.828230	9.156280	8.701544	-3.851494

```
In [96]: 1 / df
Out[96]:
```

	A	B	C	D
0	3.678365	-2.353094	1.763605	3.620145
1	-0.919624	-1.484363	8.799067	-0.676395
2	1.904807	2.470934	1.732964	-0.583090
3	-0.962215	-2.697986	-0.863638	-0.743875
4	1.183593	0.929567	-9.170108	0.608434
5	-0.680555	2.800959	-1.482360	-0.562777
6	-1.032084	-0.772485	2.416988	3.614523
7	-2.118489	-71.634509	-2.758294	-162.507295
8	-1.083352	1.116424	1.241860	-0.828904
9	0.389765	0.698687	0.746097	-0.854483

```
In [97]: df ** 4
Out[97]:
```

	A	B	C	D
0	0.005462	3.261689e-02	0.103370	5.822320e-03
1	1.398165	2.059869e-01	0.000167	4.777482e+00

(continues on next page)

(continued from previous page)

```

2    0.075962  2.682596e-02  0.110877  8.650845e+00
3    1.166571  1.887302e-02  1.797515  3.265879e+00
4    0.509555  1.339298e+00  0.000141  7.297019e+00
5    4.661717  1.624699e-02  0.207103  9.969092e+00
6    0.881334  2.808277e+00  0.029302  5.858632e-03
7    0.049647  3.797614e-08  0.017276  1.433866e-09
8    0.725974  6.437005e-01  0.420446  2.118275e+00
9   43.329821  4.196326e+00  3.227153  1.875802e+00

```

Boolean operators work as well:

```

In [98]: df1 = pd.DataFrame({"a": [1, 0, 1], "b": [0, 1, 1]}, dtype=bool)
In [99]: df2 = pd.DataFrame({"a": [0, 1, 1], "b": [1, 1, 0]}, dtype=bool)
In [100]: df1 & df2
Out[100]:
   a    b
0  False False
1  False  True
2   True  False

In [101]: df1 | df2
Out[101]:
   a    b
0  True  True
1  True  True
2  True  True

In [102]: df1 ^ df2
Out[102]:
   a    b
0  True  True
1  True  False
2  False  True

In [103]: ~df1
Out[103]:
   a    b
0  False  True
1   True  False
2  False  False

```

Transposing

To transpose, access the `T` attribute (also the `transpose` function), similar to an ndarray:

```

# only show the first 5 rows
In [104]: df[:5].T
Out[104]:
   0         1         2         3         4
A  0.271860 -1.087401  0.524988 -1.039268  0.844885
B -0.424972 -0.673690  0.404705 -0.370647  1.075770
C  0.567020  0.113648  0.577046 -1.157892 -0.109050
D  0.276232 -1.478427 -1.715002 -1.344312  1.643563

```

DataFrame interoperability with NumPy functions

Elementwise NumPy ufuncs (log, exp, sqrt, ...) and various other NumPy functions can be used with no issues on Series and DataFrame, assuming the data within are numeric:

```
In [105]: np.exp(df)
Out[105]:
```

	A	B	C	D
0	1.312403	0.653788	1.763006	1.318154
1	0.337092	0.509824	1.120358	0.227996
2	1.690438	1.498861	1.780770	0.179963
3	0.353713	0.690288	0.314148	0.260719
4	2.327710	2.932249	0.896686	5.173571
5	0.230066	1.429065	0.509360	0.169161
6	0.379495	0.274028	1.512461	1.318720
7	0.623732	0.986137	0.695904	0.993865
8	0.397301	2.449092	2.237242	0.299269
9	13.009059	4.183951	3.820223	0.310274

```
In [106]: np.asarray(df)
Out[106]:
```

```
array([[ 0.2719, -0.425 ,  0.567 ,  0.2762],
       [-1.0874, -0.6737,  0.1136, -1.4784],
       [ 0.525 ,  0.4047,  0.577 , -1.715 ],
       [-1.0393, -0.3706, -1.1579, -1.3443],
       [ 0.8449,  1.0758, -0.109 ,  1.6436],
       [-1.4694,  0.357 , -0.6746, -1.7769],
       [-0.9689, -1.2945,  0.4137,  0.2767],
       [-0.472 , -0.014 , -0.3625, -0.0062],
       [-0.9231,  0.8957,  0.8052, -1.2064],
       [ 2.5656,  1.4313,  1.3403, -1.1703]])
```

DataFrame is not intended to be a drop-in replacement for ndarray as its indexing semantics and data model are quite different in places from an n-dimensional array.

Series implements `__array_ufunc__`, which allows it to work with NumPy's [universal functions](#).

The ufunc is applied to the underlying array in a *Series*.

```
In [107]: ser = pd.Series([1, 2, 3, 4])
In [108]: np.exp(ser)
Out[108]:
```

```
0    2.718282
1    7.389056
2   20.085537
3   54.598150
dtype: float64
```

Changed in version 0.25.0: When multiple *Series* are passed to a ufunc, they are aligned before performing the operation.

Like other parts of the library, pandas will automatically align labeled inputs as part of a ufunc with multiple inputs. For example, using `numpy.remainder()` on two *Series* with differently ordered labels will align before the operation.

```
In [109]: ser1 = pd.Series([1, 2, 3], index=["a", "b", "c"])
```

(continues on next page)

(continued from previous page)

```
In [110]: ser2 = pd.Series([1, 3, 5], index=["b", "a", "c"])

In [111]: ser1
Out[111]:
a    1
b    2
c    3
dtype: int64

In [112]: ser2
Out[112]:
b    1
a    3
c    5
dtype: int64

In [113]: np.remainder(ser1, ser2)
Out[113]:
a    1
b    0
c    3
dtype: int64
```

As usual, the union of the two indices is taken, and non-overlapping values are filled with missing values.

```
In [114]: ser3 = pd.Series([2, 4, 6], index=["b", "c", "d"])

In [115]: ser3
Out[115]:
b    2
c    4
d    6
dtype: int64

In [116]: np.remainder(ser1, ser3)
Out[116]:
a    NaN
b    0.0
c    3.0
d    NaN
dtype: float64
```

When a binary ufunc is applied to a *Series* and *Index*, the Series implementation takes precedence and a Series is returned.

```
In [117]: ser = pd.Series([1, 2, 3])

In [118]: idx = pd.Index([4, 5, 6])

In [119]: np.maximum(ser, idx)
Out[119]:
0    4
1    5
2    6
dtype: int64
```

NumPy ufuncs are safe to apply to *Series* backed by non-ndarray arrays, for example `arrays.SparseArray`

(see *Sparse calculation*). If possible, the ufunc is applied without converting the underlying data to an ndarray.

Console display

Very large DataFrames will be truncated to display them in the console. You can also get a summary using `info()`. (Here I am reading a CSV version of the **baseball** dataset from the **plyr** R package):

```
In [120]: baseball = pd.read_csv("data/baseball.csv")

In [121]: print(baseball)
      id  player  year  stint team lg   g  ab  r   h  X2b  X3b  hr   rbi  sb
↪ cs  bb    so  ibb  hbp   sh  sf  gidp
0  88641  womact01  2006     2  CHN  NL  19  50  6  14    1    0    1   2.0  1.0
↪ 1.0  4   4.0  0.0  0.0  3.0  0.0  0.0
1  88643  schilcu01  2006     1  BOS  AL  31   2  0   1    0    0    0   0.0  0.0
↪ 0.0  0   1.0  0.0  0.0  0.0  0.0  0.0
..   ...      ...    ...    ...   ...   ...  ..   ...  ..   ...   ...   ..   ...   ...
↪ ...  ..      ...    ...    ...   ...   ...  ..   ...  ..   ...   ...   ..   ...   ...
98 89533  aloumo01  2007     1  NYN  NL  87 328 51 112   19    1   13  49.0  3.0
↪ 0.0 27  30.0 5.0 2.0  0.0  3.0 13.0
99 89534  alomasa02  2007     1  NYN  NL   8  22  1   3    1    0    0   0.0  0.0
↪ 0.0  0   3.0  0.0  0.0  0.0  0.0  0.0

[100 rows x 23 columns]

In [122]: baseball.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 23 columns):
#   Column  Non-Null Count  Dtype
---  ---
0    id      100 non-null    int64
1  player  100 non-null    object
2   year    100 non-null    int64
3  stint    100 non-null    int64
4   team    100 non-null    object
5    lg      100 non-null    object
6    g       100 non-null    int64
7   ab      100 non-null    int64
8    r       100 non-null    int64
9    h       100 non-null    int64
10  X2b       100 non-null    int64
11  X3b       100 non-null    int64
12  hr        100 non-null    int64
13  rbi       100 non-null    float64
14  sb        100 non-null    float64
15  cs        100 non-null    float64
16  bb        100 non-null    int64
17  so        100 non-null    float64
18  ibb       100 non-null    float64
19  hbp       100 non-null    float64
20  sh        100 non-null    float64
21  sf        100 non-null    float64
22  gidp      100 non-null    float64
dtypes: float64(9), int64(11), object(3)
memory usage: 18.1+ KB
```

However, using `to_string` will return a string representation of the DataFrame in tabular form, though it won't

always fit the console width:

```
In [123]: print(baseball.iloc[-20:, :12].to_string())
```

	id	player	year	stint	team	lg	g	ab	r	h	X2b	X3b
80	89474	finlest01	2007	1	COL	NL	43	94	9	17	3	0
81	89480	embreal01	2007	1	OAK	AL	4	0	0	0	0	0
82	89481	edmonji01	2007	1	SLN	NL	117	365	39	92	15	2
83	89482	easleda01	2007	1	NYN	NL	76	193	24	54	6	0
84	89489	delgaca01	2007	1	NYN	NL	139	538	71	139	30	0
85	89493	cormirh01	2007	1	CIN	NL	6	0	0	0	0	0
86	89494	coninje01	2007	2	NYN	NL	21	41	2	8	2	0
87	89495	coninje01	2007	1	CIN	NL	80	215	23	57	11	1
88	89497	clemero02	2007	1	NYA	AL	2	2	0	1	0	0
89	89498	claytro01	2007	2	BOS	AL	8	6	1	0	0	0
90	89499	claytro01	2007	1	TOR	AL	69	189	23	48	14	0
91	89501	cirilje01	2007	2	ARI	NL	28	40	6	8	4	0
92	89502	cirilje01	2007	1	MIN	AL	50	153	18	40	9	2
93	89521	bondsba01	2007	1	SFN	NL	126	340	75	94	14	0
94	89523	biggicr01	2007	1	HOU	NL	141	517	68	130	31	3
95	89525	benitar01	2007	2	FLO	NL	34	0	0	0	0	0
96	89526	benitar01	2007	1	SFN	NL	19	0	0	0	0	0
97	89530	ausmubr01	2007	1	HOU	NL	117	349	38	82	16	3
98	89533	aloumo01	2007	1	NYN	NL	87	328	51	112	19	1
99	89534	alomasa02	2007	1	NYN	NL	8	22	1	3	1	0

Wide DataFrames will be printed across multiple rows by default:

```
In [124]: pd.DataFrame(np.random.randn(3, 12))
```

```
Out [124]:
```

	0	1	2	3	4	5	6	7	8	9	10	11
0	-1.226825	0.769804	-1.281247	-0.727707	-0.121306	-0.097883	0.695775	0.341734	0.	↩		
↩	959726	-1.110336	-0.619976	0.149748								
1	-0.732339	0.687738	0.176444	0.403310	-0.154951	0.301624	-2.179861	-1.369849	-0.	↩		
↩	954208	1.462696	-1.743161	-0.826591								
2	-0.345352	1.314232	0.690579	0.995761	2.396780	0.014871	3.357427	-0.317441	-1.	↩		
↩	236269	0.896171	-0.487602	-0.082240								

You can change how much to print on a single row by setting the `display.width` option:

```
In [125]: pd.set_option("display.width", 40) # default is 80
```

```
In [126]: pd.DataFrame(np.random.randn(3, 12))
```

```
Out [126]:
```

	0	1	2	3	4	5	6	7	8	9	10	11
0	-2.182937	0.380396	0.084844	0.432390	1.519970	-0.493662	0.600178	0.274230	0.	↩		
↩	132885	-0.023688	2.410179	1.450520								
1	0.206053	-0.251905	-2.213588	1.063327	1.266143	0.299368	-0.863838	0.408204	-1.	↩		
↩	048089	-0.025747	-0.988387	0.094055								
2	1.262731	1.289997	0.082423	-0.055758	0.536580	-0.489682	0.369374	-0.034571	-2.	↩		
↩	484478	-0.281461	0.030711	0.109121								

You can adjust the max width of the individual columns by setting `display.max_colwidth`

```
In [127]: datafile = {
.....:     "filename": ["filename_01", "filename_02"],
.....:     "path": [
```

(continues on next page)

(continued from previous page)

```

.....:         "media/user_name/storage/folder_01/filename_01",
.....:         "media/user_name/storage/folder_02/filename_02",
.....:     ],
.....: }
.....:
In [128]: pd.set_option("display.max_colwidth", 30)

In [129]: pd.DataFrame(datafile)
Out[129]:
   filename                                path
0  filename_01  media/user_name/storage/fo...
1  filename_02  media/user_name/storage/fo...

In [130]: pd.set_option("display.max_colwidth", 100)

In [131]: pd.DataFrame(datafile)
Out[131]:
   filename                                path
0  filename_01  media/user_name/storage/folder_01/filename_01
1  filename_02  media/user_name/storage/folder_02/filename_02

```

You can also disable this feature via the `expand_frame_repr` option. This will print the table in one block.

DataFrame column attribute access and IPython completion

If a DataFrame column label is a valid Python variable name, the column can be accessed like an attribute:

```

In [132]: df = pd.DataFrame({"foo1": np.random.randn(5), "foo2": np.random.randn(5)})

In [133]: df
Out[133]:
   foo1    foo2
0  1.126203  0.781836
1 -0.977349 -1.071357
2  1.474071  0.441153
3 -0.064034  2.353925
4 -1.282782  0.583787

In [134]: df.foo1
Out[134]:
0    1.126203
1   -0.977349
2    1.474071
3   -0.064034
4   -1.282782
Name: foo1, dtype: float64

```

The columns are also connected to the IPython completion mechanism so they can be tab-completed:

```

In [5]: df.foo<TAB>  # noqa: E225, E999
df.foo1  df.foo2

```